Vast.ai

Peer-to-Peer GPU Platform

Intending to distribute decentralized computing to all.

PEER-TO-PEER

DECENTRALIZED

GPU POWER

DREAMSERVER INTEGRATION
https://github.com/Light-Heart-Labs/DreamServer

vast.ai Developers Pricing Products Hosting Use Cases Company Contact Sales Console

SOC 2 Certified

Create a New Account

Email

Username (Defaults to email)

Password

Password (Repeat)

I have read the Terms of Use and Privacy Policy I have read and agree to the terms of the Terms and Conditions. I consent to the use and processing of my personal information as set forth therein and I always agree to use Vast. AI's site and services in accordance therewith.

Subscribe to Vast marketing emails

Sign Up

Continue with Google

Continue with Github

Or Sign in • Reset password

Earth Auto Sort

verified Max Duration 27 days $0.457/hr plus bandwidth Reliability 99.10% RENT

verified Max Duration 69 mon. $0.623/hr plus bandwidth Reliability 99.46% RENT

verified Max Duration 14 days $1.700/hr plus bandwidth Reliability 99.955% RENT

verified Max Duration 10 days $1.082/hr plus bandwidth Reliability 99.77% RENT

verified Max Duration 3 mon. $0.129/hr plus bandwidth Reliability 99.72% RENT

Questions? Suggestions? Talk to us!

M morissonla... Individual

Search
Templates
Instances
Storage
Serverless NEW

Account
Billing
Earnings
Members
Audit Logs
Keys
Settings

Credits

Remaining balance

Add Credit

Current Usage (excluding bandwidth)

Total:	0.423 $/hr	10.152 $/day
GPU:	0.36 $/hr	8.64 $/day
Disk:	0.063 $/hr	1.512 $/day
Volumes:	0.000 $/hr	0.000 $/day

You will run out of credits in ~6 hours, 56 minutes

Notifications

Low balance notifications

Get an email notification when your balance goes below your credit threshold

Credit threshold

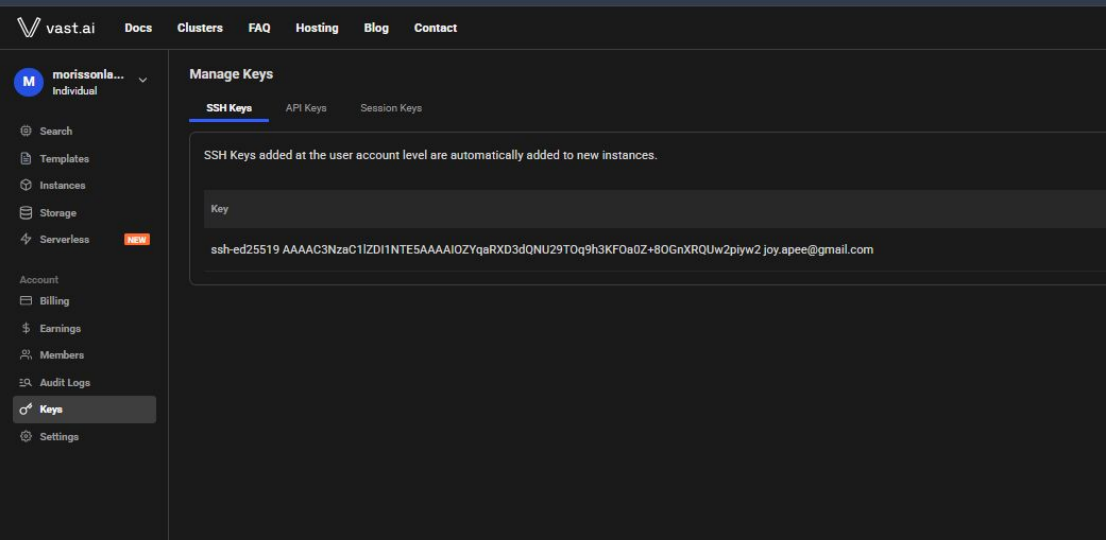

Linux / macOS (Terminal)

Generate SSH key pair:
`ssh-keygen -t ed25519 -C your_email@example.com`
Press Enter for default location (`~/.ssh/id_ed25519`).
Optionally set a passphrase.
Creates: Private key (`~/.ssh/id_ed25519`) and Public key (`~/.ssh/id_ed25519.pub`).
Show public key:
`cat ~/.ssh/id_ed25519.pub`

Press on key, then paste the public key.

SSH key pairs are crucial for secure authentication and establishing encrypted connections between computers.

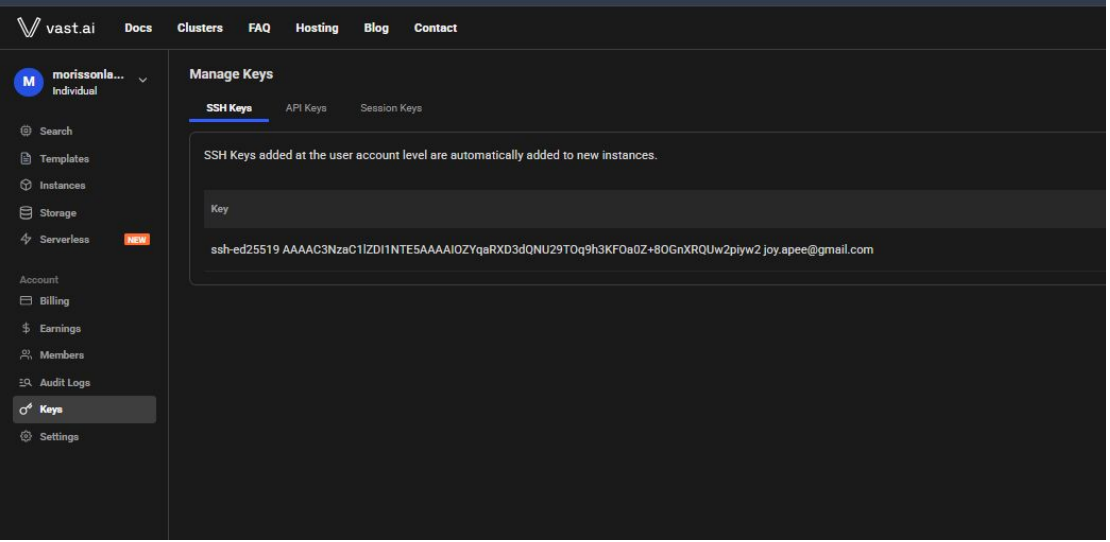

Press on key, then paste the public key.

SSH key pairs are crucial for secure authentication and establishing encrypted connections between computers.

Windows (PowerShell or Git Bash)

Generate SSH key pair:
`ssh-keygen -t ed25519 -C "your_email@example.com"`
Press Enter for default location (`C:\Users\YourUser\.ssh\id_ed25519` or `~/.ssh/id_ed25519` in Git Bash).
Optionally set a passphrase.
Creates: Private key (`id_ed25519`) and Public key (`id_ed25519.pub`).
Show public key:
`cat ~/.ssh/id_ed25519.pub` (Use this in Git Bash)
`Get-Content $HOME\.ssh\id_ed25519.pub` (Use this in PowerShell)

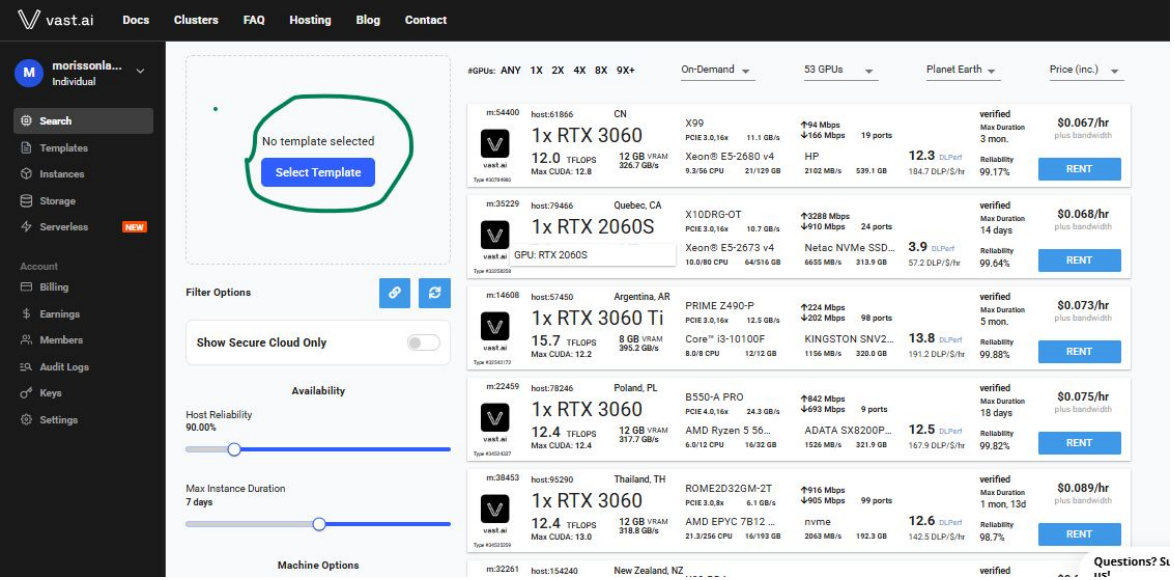

To begin, create an account with a click and then recharge at least $5. The platform supports all payment methods and does not store card information.

Once the account is ready, we can focus on a simple setup with just a few steps.

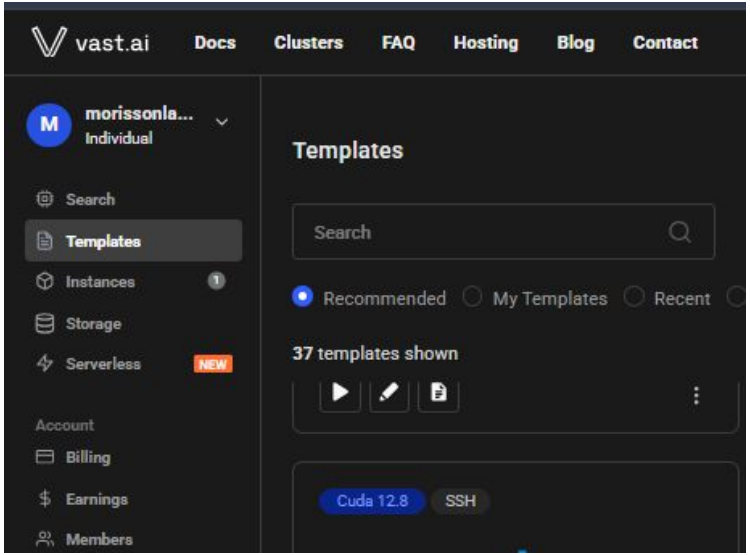

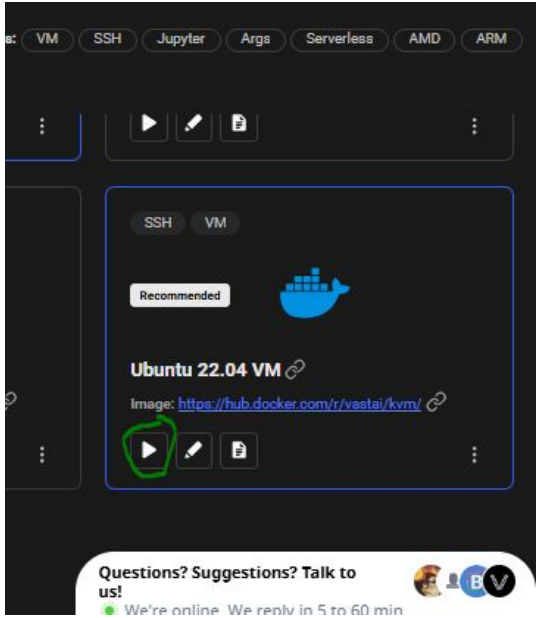

Press on template then select Ubuntu 22.04 VM

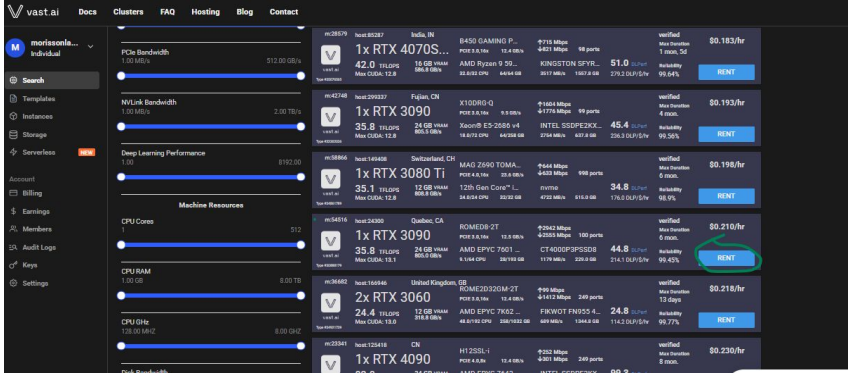

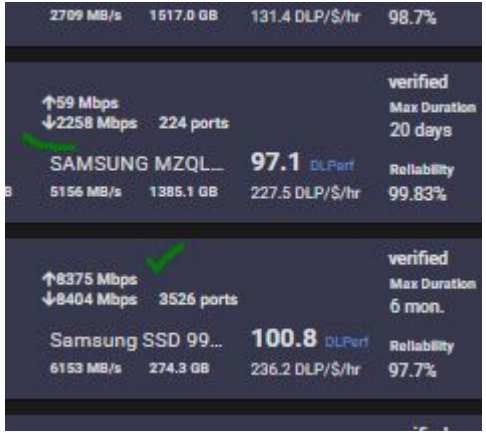

Consider renting a GPU with high network speed to ensure quicker installation times.

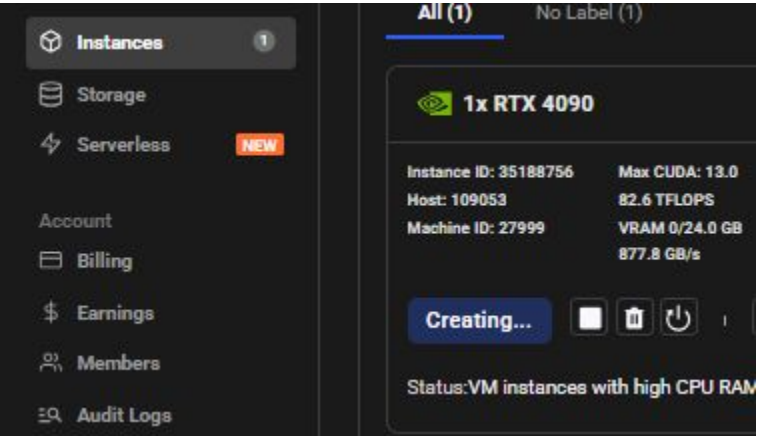

Select "connect" to obtain the IP address and port number.

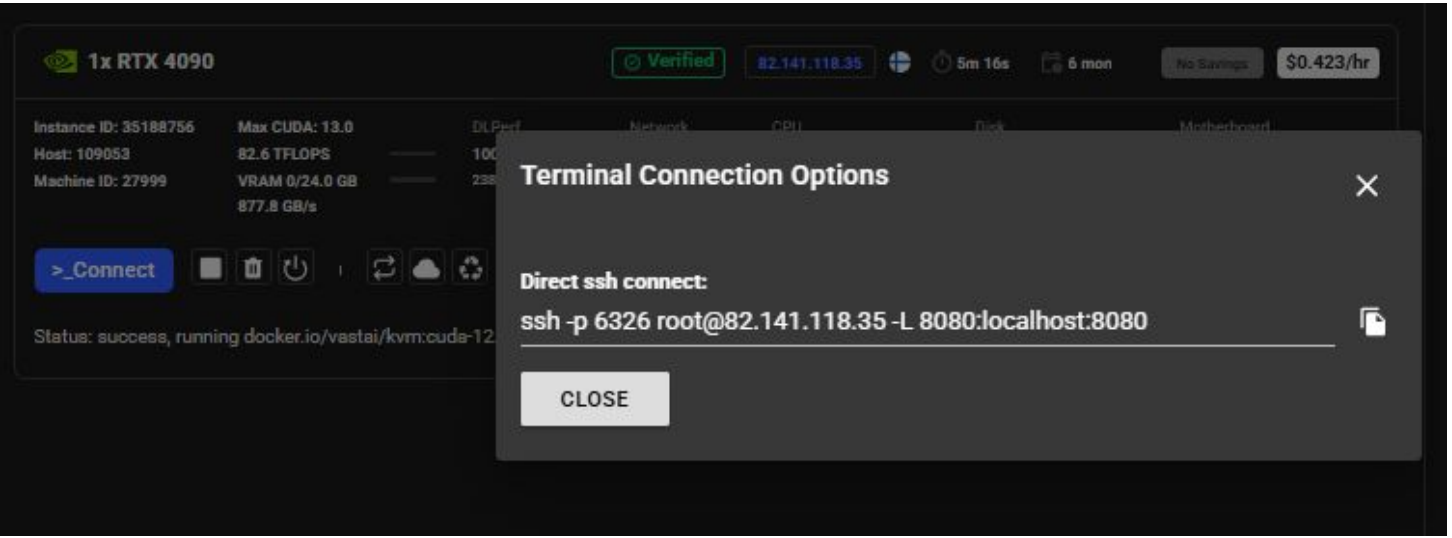

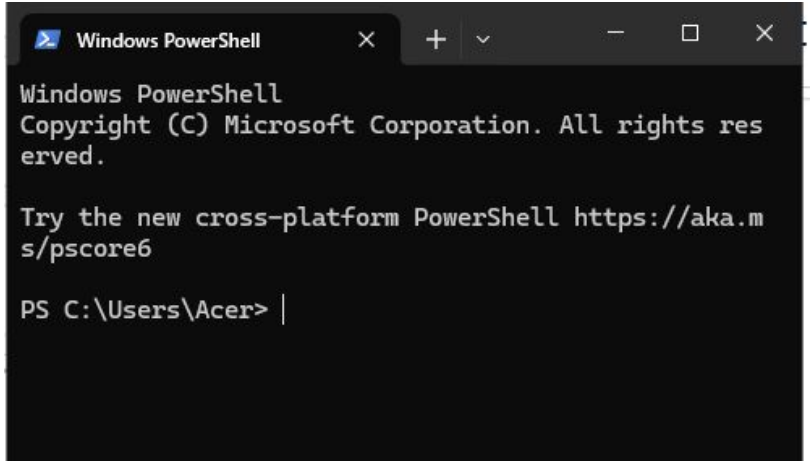

Establish a secure shell connection, utilizing the specified identity file, port, and local port forwarding configuration:

```
ssh -i "C:/Users/Acer/.ssh/id_ed25519" -p
6326 -L 8080:localhost:8080
root@82.141.118.35
```

Note: Ensure the identity profile path is replaced with the correct local file path.

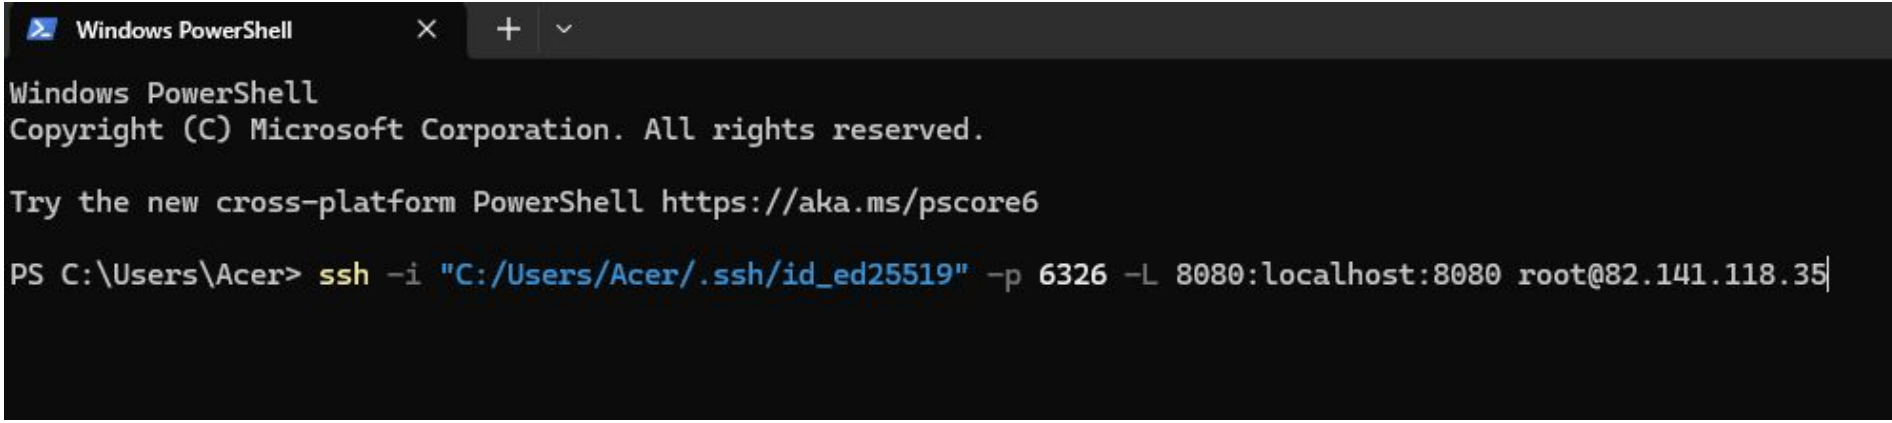

```
PS C:\Users\Acer> ssh -i "C:/Users/Acer/.ssh/id_ed25519" -p 6326 -L 8080:localhost:8080 root@82.141.118.35
The authenticity of host '[82.141.118.35]:6326 ([82.141.118.35]:6326)' can't be established.
ED25519 key fingerprint is SHA256:/C2/[redacted]mcM.
This host key is known by the following other names/addresses:
    C:\Users\Acer/.ssh/known_hosts:929: [94.63.174.45]:14195
    C:\Users\Acer/.ssh/known_hosts:932: [142.115.221.60]:51623
    C:\Users\Acer/.ssh/known_hosts:933: [96.38.133.243]:16551
    C:\Users\Acer/.ssh/known_hosts:937: [155.103.252.98]:4851
    C:\Users\Acer/.ssh/known_hosts:942: [69.157.137.231]:16343
    C:\Users\Acer/.ssh/known_hosts:946: [142.115.221.60]:51602
    C:\Users\Acer/.ssh/known_hosts:947: [159.48.242.1]:35367
    C:\Users\Acer/.ssh/known_hosts:948: [142.115.221.60]:51562
    (7 additional names omitted)
Are you sure you want to continue connecting (yes/no/[fingerprint])? yes
```

cd /home/ubuntu/

git clone -b feat/auto-adaptive-vastai-setup https://github.com/Arifuzzamanjoy/DreamServer.git

cd DreamServer/resources/p2p-gpu/

./setup.sh

```
root@ubuntu:/home/ubuntu# git clone -b feat/auto-adaptive-vastai-setup https://github.com/Arifuzzamanjoy/DreamServer.git
Cloning into 'DreamServer'...
remote: Enumerating objects: 13790, done.
remote: Counting objects: 100% (4625/4625), done.
remote: Compressing objects: 100% (1049/1049), done.
remote: Total 13790 (delta 3993), reused 3616 (delta 3575), pack-reused 9165 (from 3)
Receiving objects: 100% (13790/13790), 17.36 MiB | 34.25 MiB/s, done.
Resolving deltas: 100% (8899/8899), done.
root@ubuntu:/home/ubuntu# cd DreamServer/resources/p2p-gpu/
root@ubuntu:/home/ubuntu/DreamServer/resources/p2p-gpu#
root@ubuntu:/home/ubuntu/DreamServer/resources/p2p-gpu#
root@ubuntu:/home/ubuntu/DreamServer/resources/p2p-gpu# ./setup.sh
```

```
root@ubuntu:/home/ubuntu/DreamServer/resources/p2p-gpu# ./setup.sh

  DreamServer - Vast.ai Setup v6.1.0
  https://github.com/Light-Heart-Labs/DreamServer


── Phase 0/12: Preflight checks ──

[✓] NVIDIA GPU: NVIDIA GeForce RTX 4090 × 1 (24564 MiB VRAM each)
[✓] GPU backend: nvidia | CPUs: 27 | Disk: 157 GB
[✓] Docker Compose: docker compose (2.35.1)
[✓] Verifying NVIDIA Docker passthrough (timeout 180s; first run may pull nvidia/cuda:12.4.1-base-ubuntu22.04)
[!] NVIDIA GPU passthrough test failed (exit 125) - checking toolkit...
[!] nvidia-container-toolkit not installed - attempting install
File '/usr/share/keyrings/nvidia-container-toolkit-keyring.gpg' exists. Overwrite? (y/N)
```

Execute these commands exactly as shown and press Enter.
Select Y (recommended)
After this step, take a break while your dream server is being prepared.

```
── SSH Tunnel (recommended for Vast.ai) ──

  Windows PowerShell (all ports, recommended):
ssh -N -o ExitOnForwardFailure=yes -p 6326 -i $env:USERPROFILE\.ssh\id_ed25519 -L 58080:127.0.0.1:3001 -L 11434:127.0.0.1:11434 -L 3000:127.0.0.1:30
00 -L 8888:127.0.0.1:8888 -L 3004:127.0.0.1:3004 -L 9000:127.0.0.1:9000 -L 8880:127.0.0.1:8880 -L 6333:127.0.0.1:6333 -L 6334:127.0.0.1:6334 -L 8090
:127.0.0.1:8090 -L 4000:127.0.0.1:4000 -L 7860:127.0.0.1:7860 -L 3006:127.0.0.1:3006 -L 5678:127.0.0.1:5678 -L 8085:127.0.0.1:8085 -L 7890:127.0.0.1
:7890 -L 3010:127.0.0.1:3010 -L 3003:127.0.0.1:3003 -L 8188:127.0.0.1:8188 -L 3001:127.0.0.1:3001 -L 3005:127.0.0.1:3005 -L 3002:127.0.0.1:3002 root
@82.141.118.35
Open dashboard: http://127.0.0.1:3001/
Easy alias:    http://127.0.0.1:58080/

  Linux / macOS (all ports):
ssh -N -p 6326 -i ~/.ssh/id_ed25519 -L 58080:127.0.0.1:3001 -L 11434:127.0.0.1:11434 -L 3000:127.0.0.1:3000 -L 8888:127.0.0.1:8888 -L 3004:127.0.0.1
:3004 -L 9000:127.0.0.1:9000 -L 8880:127.0.0.1:8880 -L 6333:127.0.0.1:6333 -L 6334:127.0.0.1:6334 -L 8090:127.0.0.1:8090 -L 4000:127.0.0.1:4000 -L 7
860:127.0.0.1:7860 -L 3006:127.0.0.1:3006 -L 5678:127.0.0.1:5678 -L 8085:127.0.0.1:8085 -L 7890:127.0.0.1:7890 -L 3010:127.0.0.1:3010 -L 3003:127.0.
0.1:3003 -L 8188:127.0.0.1:8188 -L 3001:127.0.0.1:3001 -L 3005:127.0.0.1:3005 -L 3002:127.0.0.1:3002 root@82.141.118.35
Open dashboard: http://127.0.0.1:3001/

  Auto-reconnect scripts:
  Windows: scp -P 6326 root@82.141.118.35:/home/dream/dream-server/connect-tunnel.ps1 .
           powershell -ExecutionPolicy Bypass -File .  Linux/macOS/WSL: scp -P 6326 root@82.141.118.35:/home/dream/dream-server/connect-tunnel.sh .
                   FULL_TUNNEL=1 bash connect-tunnel.sh

  If Windows reports "bind [127.0.0.1]:PORT: Permission denied",
  that local port is reserved by Hyper-V/WinNAT. Use a different local port:
    -L 58080:127.0.0.1:3001   (or any free high port)
  Optional admin fix:
    net stop winnat; net start winnat   (run PowerShell as admin)
  Check excluded ranges: netsh int ipv4 show excludedportrange protocol=tcp
  If you see "channel N: open failed: connect failed: Connection refused",
  the SSH tunnel is up, but that specific remote service is not listening yet.
```

ssh -N -o ExitOnForwardFailure=yes -p 6326 -i C:/Users/Acer/.ssh/id_ed25519 -L 58080:127.0.0.1:3001 -L 11434:127.0.0.1:11434 -L 3000:127.0.0.1:3000 -L 8888:127.0.0.1:8888 -L 3004:127.0.0.1:3004 -L 9000:127.0.0.1:9000 -L 8880:127.0.0.1:8880 -L 6333:127.0.0.1:6333 -L 6334:127.0.0.1:6334 -L 8090:127.0.0.1:8090 -L 4000:127.0.0.1:4000 -L 7860:127.0.0.1:7860 -L 5678:127.0.0.1:5678 -L 8085:127.0.0.1:8085 -L 7890:127.0.0.1:7890 -L 3010:127.0.0.1:3010 -L 8188:127.0.0.1:8188 -L 3001:127.0.0.1:3001 -L 3005:127.0.0.1:3005 -L 3002:127.0.0.1:3002 root@82.141.118.35

```
PS C:\Users\Acer>
PS C:\Users\Acer> ssh -N -o ExitOnForwardFailure=yes -L 58080:127.0.0.1:3001 -L 11434:127.0.0.1:11434 -L 3000:127.0.0.1:
3000 -L 8888:127.0.0.1:8888 -L 3004:127.0.0.1:3004 -L 9000:127.0.0.1:9000 -L 8880:127.0.0.1:8880 -L 6333:127.0.0.1:6333
-L 6334:127.0.0.1:6334 -L 8090:127.0.0.1:8090 -L 4000:127.0.0.1:4000 -L 7860:127.0.0.1:7860 -L 5678:127.0.0.1:5678 -L 80
85:127.0.0.1:8085 -L 7890:127.0.0.1:7890 -L 3010:127.0.0.1:3010 -L 8188:127.0.0.1:8188 -L 3001:127.0.0.1:3001 -L 3005:12
7.0.0.1:3005 -L 3002:127.0.0.1:3002 p2p_gpu_all
```

```
PS C:\Users\Acer> ssh -N -o ExitOnForwardFailure=yes -p 6326 -i C:/Users/Acer/.ssh/id_ed25519 -L 58080:127.0.0.1:3001 -L
 11434:127.0.0.1:11434 -L 3000:127.0.0.1:3000 -L 8888:127.0.0.1:8888 -L 3004:127.0.0.1:3004 -L 9000:127.0.0.1:9000 -L 88
80:127.0.0.1:8880 -L 6333:127.0.0.1:6333 -L 6334:127.0.0.1:6334 -L 8090:127.0.0.1:8090 -L 4000:127.0.0.1:4000 -L 7860:12
7.0.0.1:7860 -L 5678:127.0.0.1:5678 -L 8085:127.0.0.1:8085 -L 7890:127.0.0.1:7890 -L 3010:127.0.0.1:3010 -L 8188:127.0.0
.1:8188 -L 3001:127.0.0.1:3001 -L 3005:127.0.0.1:3005 -L 3002:127.0.0.1:3002 root@82.141.118.35
```

recommended

After installation open a fresh terminal /powershell (administrator privilege optional)

Replace the current IP port and identity profile with your specific configurations. Utilize your own Dreamserver for only $5, requiring no GPU or initial investment.

```
Services:
  llama-server (LLM Inference): http://localhost:11434
  Open WebUI (Chat):     http://localhost:3000
  SearXNG (Web Search):  http://localhost:8888
  Perplexica (Deep Research): http://localhost:3004
  Whisper (STT):         http://localhost:9000
  Kokoro (TTS):          http://localhost:8880
  Qdrant (Vector DB):    http://localhost:6333
  QDRANT_GRPC_PORT:      http://localhost:6334
  TEI (Embeddings):      http://localhost:8090
  LiteLLM (API Gateway): http://localhost:4000
  OpenClaw (Agents):     http://localhost:7860
  Langfuse (LLM Observability): http://localhost:3006
  n8n (Workflows):       http://localhost:5678
  Privacy Shield (PII Protection): http://localhost:8085
  APE (Agent Policy Engine): http://localhost:7890
  DreamForge:            http://localhost:3010
  OpenCode (IDE):        http://localhost:3003
  ComfyUI (Image Generation): http://localhost:8188
  Dashboard (Control Center): http://localhost:3001
  Token Spy (Usage Monitor): http://localhost:3005
  Dashboard API (System Status): http://localhost:3002

Upload Custom Models:
  scp -P 6326 my-model.gguf root@82.141.118.35:/home/dream/dream-server/data/models/
  # Then: edit .env GGUF_FILE=my-model.gguf && docker restart dream-llama-server

Commands:
  bash setup.sh --status     # Check health
  bash setup.sh --info       # Connection details
  bash setup.sh --fix        # Apply fixes + restart
  bash setup.sh --resume     # Quick restart after SSH drop
  bash setup.sh --teardown   # Stop all (save billing)
```

```
root@ubuntu: ~        Administrator: Windows Powe

Windows PowerShell
Copyright (C) Microsoft Corporation. All rights reserved.

Try the new cross-platform PowerShell https://aka.ms/pscore6

PS C:\Users\Acer> ssh -N -o ExitOnForwardFailure=yes -L 58080:127.0.0.1:3001 -L 11434:127.0.0.1:11434 -L 3000:127.0.0.1:
3000 -L 8888:127.0.0.1:8888 -L 3004:127.0.0.1:3004 -L 9000:127.0.0.1:9000 -L 8880:127.0.0.1:8880 -L 6333:127.0.0.1:6333
-L 6334:127.0.0.1:6334 -L 8090:127.0.0.1:8090 -L 4000:127.0.0.1:4000 -L 7860:127.0.0.1:7860 -L 5678:127.0.0.1:5678 -L 80
85:127.0.0.1:8085 -L 7890:127.0.0.1:7890 -L 3010:127.0.0.1:3010 -L 8188:127.0.0.1:8188 -L 3001:127.0.0.1:3001 -L 3005:12
7.0.0.1:3005 -L 3002:127.0.0.1:3002 p2p_gpu_all
```

```
root@ubuntu: ~        Administrator: Windows Power

root@ubuntu:~#
root@ubuntu:~#
```

```
root@ubuntu:/home/ubuntu/DreamServer/resources/p2p-gpu#
setup.sh --fix
```

If get issue

Just paste this on vm terminal

localhost:3001

Dream Server

LOCAL AI // SOVEREIGN INTELLIGENCE

Standard • v2.0.0

- Dashboard
- Extensions
- Models
- Settings

QUICK LINKS VIEW ALL

- APE (Agent Policy Engine) OPEN
- ComfyUI (Image Generation) OPEN
- Dashboard (Control Center) OPEN
- DreamForge OPEN
- TEI (Embeddings) OPEN
- LiteLLM (API Gateway) OPEN

Services Online: 19/19

VRAM 21.2/24 GB

Dashboard

19/19 services online.

Standard qwen3-30b-a3b v2.0.0

- AI Chat • READY
- Voice Assistant • READY
- Document Q&A • READY
- Workflow Automation • READY
- Image Generation • READY
- Coding Assistant • READY
- Agent Governance • READY
- Agentic Coding • READY
- RAG Pipeline • READY

System Status

SIGNAL GRAPH

Inference telemetry

CONTEXT 32K

GPU TEMP 29°C | MEM FREE 35.1 GB

Tokens per second

LIVE THROUGHPUT 0

13 7 0 8.0

Tokens generated

ACCUMULATED OUTPUT 0

1.3k 672 0 1.2k

GPU 0% GeForce RTX 4090

VRAM 21.2 GB of 24 GB

CPU 0.3% utilization

RAM 8.8 GB of 43.9 GB

GPU POWER 5.8W live

GPU TEMP 29°C normal

CONTEXT 32k max tokens

UPTIME 1h 18m system

MODEL Qwen3-30B-A3B-Q4_K_M...